

# MAP REDUCE HADOOP (STREAMING AND CHAINING JOBS)

---

Master Program in Computer Science  
University of Calabria

*Prof. F. Ricca*

# Streaming in Hadoop

- One can use Java to write Hadoop programs
- Hadoop supports other languages
  - Via a generic API called **Streaming**
  - Useful for writing simple, short MapReduce scripts
  - Any scripting/programming language using non-Java libraries
- Hadoop Streaming
  - Interacts with programs using the Unix streaming paradigm
  - Inputs come in through STDIN
  - Outputs go to STDOUT
  - *Data has to be text based and each line is considered a record*

# Streaming Data Flow

- Similar to the following Unix command line:

```
cat [input_file] | [mapper] | sort  
                        | [reducer] >[output_file]
```

- Only need to invoke the Streaming jar

```
bin/hadoop jar [hadoop-streaming-X.Y.Z.jar]  
    -input      [input_file]  
    -output     [output_files]  
    -file       [script to share]  
    -mapper     'mappercommand'  
    -reducer    'reducercommand'  
    ...
```

# Word count per split

```
bin/hadoop jar hadoop-streaming-X.Y.Z.jar  
    -input    input.txt  
    -output   count  
    -mapper   'wc -l'  
    -numReduceTasks 0
```

# Working principles

- Works with any executable script
  - Reads a line-oriented data stream from STDIN
  - Outputs to STDOUT with Hadoop Streaming
- Each mapper
  - Sees the entire stream of data
  - Takes on the responsibility of breaking the stream into (line-oriented) record
  - In Java the framework itself breaks input data
- Streaming works on key/value pairs just like in Java
  - The tab character to separate the key from the value in a record
  - When there's no tab character, the entire record is considered the key and the value is empty text

# Streaming with key/value pairs

1. The mapper
  - Reads a split through STDIN
  - Extracts each line as a record  
(can interpret each input record as a key/value pair or a line of text)
2. The framework ensures that
  - Each line of your mapper's output is a key/ value pair
  - Key/value pairs with the same key will end up at the same reducer
3. The reducer
  - key/value pairs are sorted by key
  - is responsible for performing the grouping,
  - reads one line at a time from STDIN and keeps track of the keys
4. The reducer output (STDOUT) is written to a file directly

# Implement Max via script

```
./bin/hadoop jar hadoop-streaming-X.Y.Z.jar  
  -input input.txt  
  -output out.txt  
  -mapper 'cut -f1,3'  
  -reducer 'maxGroupAttr.py'
```

# Equivalent to

```
cat input.txt | cut -f1,3  
  | sort | ./maxGroupAttr.py
```

# maxGroupAttr.py

```
#!/usr/bin/python
import sys

lastKey=None
max    = 0
for line in sys.stdin:
    (key, val) = line.strip().split()
    if lastKey and lastKey!=key:
        print lastKey + "\t" + max
        max=0
        lastKey=None
    lastKey=key
    if (val > max):
        max = val
print lastKey + "\t" + max
```



# Streaming with the Aggregate package

- The library package called Aggregate
  - Simplifies obtaining aggregate statistics of a data set
  - Can simplify the writing of Java statistics collectors
    - Especially when used with Streaming
- Each line of the mapper's output looks like

```
function:key\tvalue
```

- The Aggregate reducer applies the function to the set of values for each key

# Standard aggregation functions

Value aggregator	Description
DoubleValueSum	Sums up a sequence of double values.
LongValueMax	Finds the maximum of a sequence of long values.
LongValueMin	Finds the minimum of a sequence of long values.
LongValueSum	Sums up a sequence of long values.
StringValueMax	Finds the lexicographical maximum of a sequence of string values.
StringValueMin	Finds the lexicographical minimum of a sequence of string values.
UniqValueCount	Finds the number of unique values (for each key).
ValueHistogram	Finds the count, minimum, median, maximum, average, and standard deviation of each value. (See text for further explanation.)

# Generic aggregator example

```
./bin/hadoop jar hadoop-streaming-X.Y.Z.jar  
-input input.txt  
-output out.txt  
-mapper 'funtionMapper.py LongValueSum 0'  
-reducer aggregate
```

```
./bin/hadoop jar hadoop-streaming-X.Y.Z.jar  
-input input.txt  
-output out.txt  
-mapper 'funtionMapper.py DoubleValueSum 0 2'  
-reducer aggregate
```

# functionMapper.py

```
#!/usr/bin/python
import sys

function = (sys.argv[1])
keyIdx = int(sys.argv[2])

if len(sys.argv) > 3:
    valIdx = int(sys.argv[3])
else:
    valIdx = -1

for line in sys.stdin:
    fields = line.strip().split()
    if fields[keyIdx] and not fields[keyIdx].startswith("#"):
        if valIdx > 0:
            print function + ":" + fields[keyIdx] + "\t" + fields[valIdx]
        else:
            print function + ":" + fields[keyIdx] + "\t" + "1"
```

**(END)**

Usage: \$HADOOP\_HOME/bin/hadoop jar hadoop-streaming.jar [options]

Options:

-input	<path> DFS input file(s) for the Map step.
-output	<path> DFS output directory for the Reduce step.
-mapper	<cmd JavaClassName> Optional. Command to be run as mapper.
-combiner	<cmd JavaClassName> Optional. Command to be run as combiner.
-reducer	<cmd JavaClassName> Optional. Command to be run as reducer.
-file	<file> Optional. File/dir to be shipped in the Job jar file. Deprecated. Use generic option "-files" instead.
-inputformat	<TextInputFormat(default) SequenceFileAsTextInputFormat JavaClassName> Optional. The input format class.
-outputformat	<TextOutputFormat(default) JavaClassName> Optional. The output format class.
-partitioner	<JavaClassName> Optional. The partitioner class.
-numReduceTasks	<num> Optional. Number of reduce tasks.
-inputreader	<spec> Optional. Input recordreader spec.
-cmdenv	<n>=<v> Optional. Pass env.var to streaming commands.
-mapdebug	<cmd> Optional. To run this script when a map task fails.
-reduceddebug	<cmd> Optional. To run this script when a reduce task fails.
-io	<identifier> Optional. Format to use for input to and output from mapper/reducer commands
-lazyOutput	Optional. Lazily create Output.
-background	Optional. Submit the job and don't wait till it completes.
-verbose	Optional. Print verbose output.
-info	Optional. Print detailed usage.
-help	Optional. Print help message.

# Chaining MapReduce jobs (1)

- Many complex tasks need to be broken down
  - Each accomplished by an individual MapReduce job

- Chaining MapReduce jobs in a sequence

- Like:

`mapreduce-1 | mapreduce-2 | ...`

- Just schedule every next job after the other in the main

`Job.waitForCompletion(true) ? 0 : 1`

- To wait until preceding job completes!!

# Chaining MapReduce jobs (2)

- Complex chains using the **ChainMapper** class

- Like:

`[map]+ | reduce [| map]*`

- `pre-processing | standard computation | post-processing`

...

```
ChainMapper.addMapper(job,map1.class,key.class,val.class)
```

```
ChainMapper.addMapper(job,map2.class,key.class,val.class)
```

```
ChainMapper.setReducer(job,red.class,key.class,val.class)
```

```
ChainMapper.addMapper(job,map3.class,key.class,val.class)
```

...

```
Job.waitForCompletion(true) ? 0 : 1
```

# Chaining MapReduce jobs (3)

- Enforce specific dependencies
  - Specify like a chain of dependencies
  - Using **JobControl** and **ControlledJob** classes

...

```
ControlledJob cJ1 = new ControlledJob(job1)
ControlledJob cJ2 = new ControlledJob(job2)
cJ1.addDependingJob(cJ2)
JobControl jC = new JobControl("A chain")
jC.addJob(cJ1)
jC.addJob(cJ1)
jC.run()
```



# LET'S START

---

(open, configure and run examples)

# Required Software

- Eclipse
  - [www.eclipse.org](http://www.eclipse.org)
- Hadoop
  - Download it from <http://hadoop.apache.org/>
  - Obtain everything it via Maven  
(<https://mvnrepository.com/artifact/org.apache.hadoop>)
- Small examples do not need a cluster
  - **Running Hadoop in Standalone mode**
- Terminal app
- Few python scripts
- Few data files