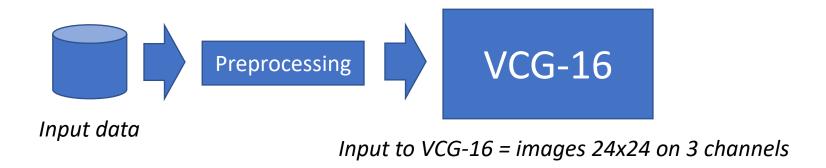




- ▶ The VCG-16 network is a Deep ConvNet with 16 layers
 - ▶ The model has been trained on the ImageNet ILSVRC-2012 dataset, which includes images of 1,000 classes, and is split into three sets: training (1.3 million images), validation (50,000 images), and testing (100,000 images). Each image is (224×224) on 3 channels.
- The network (as well as others) is available in Keras
 - ▶ Building image-recognition application is rather easy

Tre-trained Networks

- ▶ VCG-16, as well as many other networks, is available in Keras
 - ▶ Re-using these networks requires some pre-processing on the input data



2023/2024 **MACHINE LEARNING**

VCG-16 Documentation

```
tf.keras.applications.VGG16(
    include_top=True,
    weights="imagenet",
    input_tensor=None,
    input_shape=None,
    pooling=None,
    classes=1000,
    classifier_activation="softmax",
```

Arguments

- **include_top**: whether to include the 3 fully-connected layers at the top of the network.
- weights: one of None (random initialization), 'imagenet' (pre-training on ImageNet), or the path to the weights file to be loaded.
- **input_tensor**: optional Keras tensor (i.e. output of layers.Input()) to use as image input for the model.
- **input_shape**: optional shape tuple, only to be specified if include_top is False (otherwise the input shape has to be (224, 224, 3) (with channels_last data format) or (3, 224, (with channels_first data format). It should have exactly 3 input channels, and width and height should be no smaller than 32. E.g. (200, 200, 3) would be one valid value.
- pooling: Optional pooling mode for feature extraction when include top is False. -None means that the output of the model will be the 4D tensor output of the last convolutional block. - avg means that global average pooling will be applied to the output of the last convolutional block, and thus the output of the model will be a 2D tensor. - max means that global max pooling will be applied.
- **classes**: optional number of classes to classify images into, only to be specified if include_top is True, and if no weights argument is specified.
- **classifier_activation**: A str or callable. The activation function to use on the "top" layer. Ignored unless include top=True. Set classifier_activation=None to return the logits of the "top" layer.

Example (1/4)

```
import tensorflow as tf
from tensorflow.keras.applications.vgg16 import VGG16
import matplotlib.pyplot as plt
import numpy as np
import cv2

# prebuild model with pre-trained weights on imagenet
model = VGG16(weights='imagenet', include_top=True)
model.compile(optimizer='sgd', loss='categorical_crossentropy')
```

2023/2024

Example (2/4)

summarize the mdodel
model.summary()

Layer (type)	Output Shape	Param #
input_4 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168

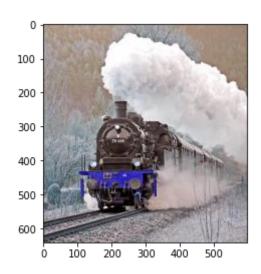
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
fc1 (Dense)	(None, 4096)	102764544
fc2 (Dense)	(None, 4096)	16781312
predictions (Dense)	(None, 1000)	4097000

Example (3/4)

```
img = cv2.imread('steam-locomotive.jpg')
print(img.shape)
plt.imshow(img)
```

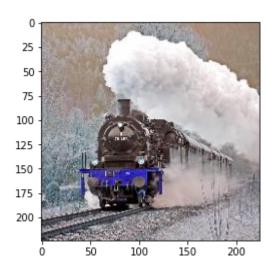


(640, 598, 3)



```
# resize into VGG16 trained images' format
# add a dummy iniziatl axis
im = cv2.resize(img, (224, 224))
plt.imshow(im)
im = np.expand_dims(im, axis=0)
im.astype(np.float32)
print(im.shape)
```





Example (4/4)

```
# predict
out = model.predict(im)
index = np.argmax(out)
print(index)
```

820

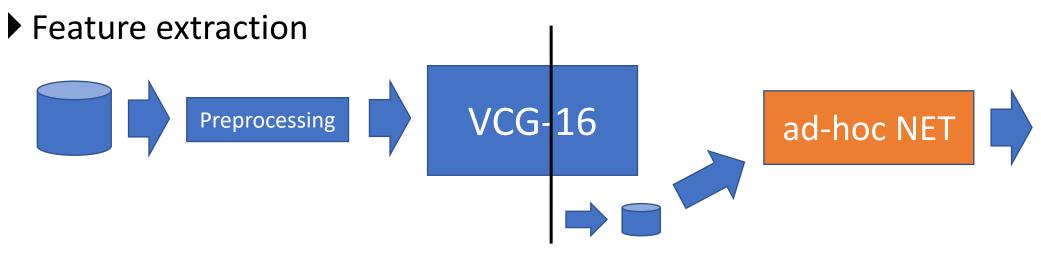
Transfer learning

- ▶ The knowledge inferred for solving a task can be reused in other tasks
- ▶ A VCG-16 network can be used to solve different classification problems
 - ▶ But we need to focus on the first layers, which extract the relevant features
 - Indeed, deeper layers are too specific for the application at hand

2023/2024 **MACHINE LEARNING**

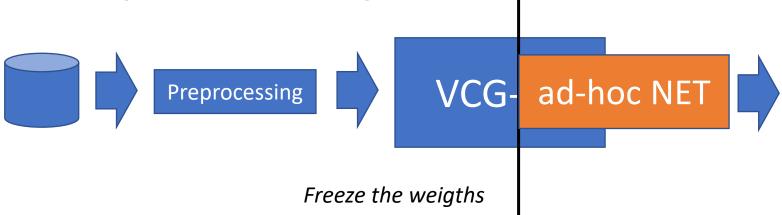


Usage of pre-trained nets

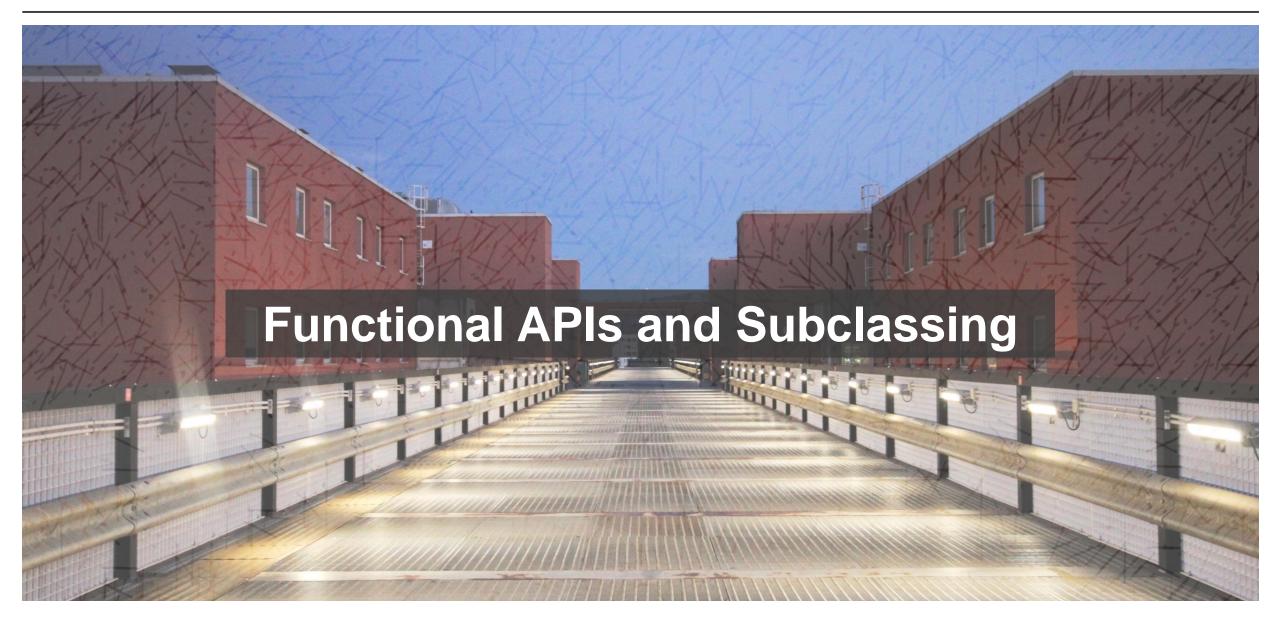


Features extracted from some given layer



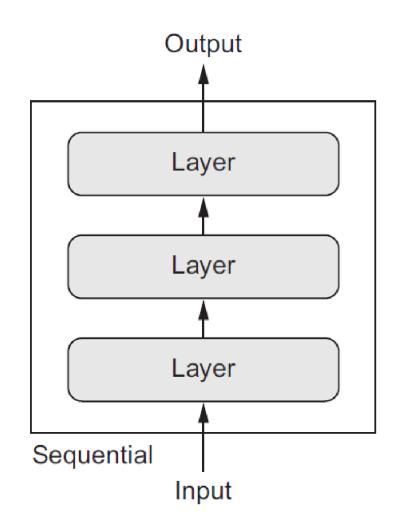






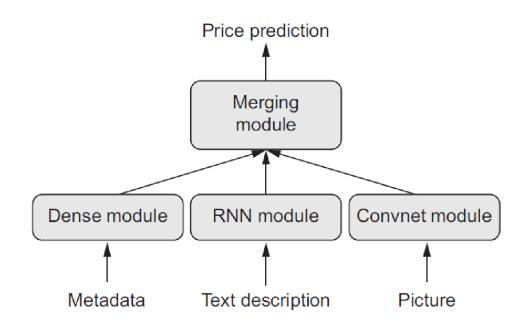
Limits of sequential models

- There are a number of applications where information does not just flow from an input to an output
- The sequential model is not flexible

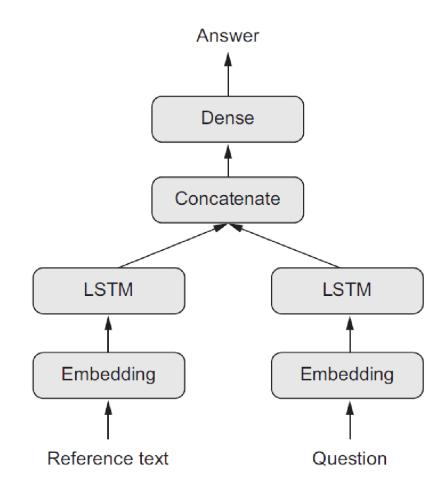


Multi-input

I The various inputs need different kinds of specialized architectures to be dealt with



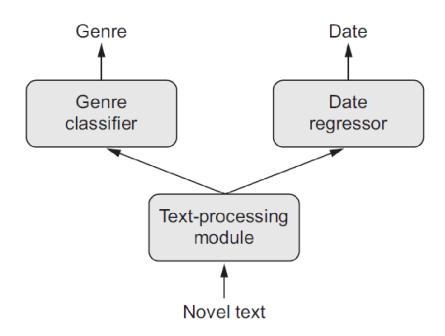
Price prediction



Question answering

Multi-output

There are applications in which we predict different kinds of heterogenous information



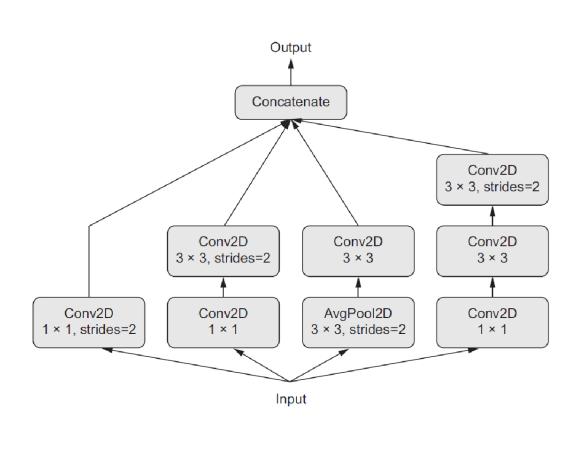
Age Gender Income Dense Dense Dense 1D convnet Social media posts

Advanced novel text classification

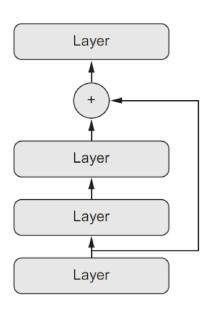
Social media analysis



Advanced Architectures



Inception



Residual

Advanced Keras - Subclassing

- Custom Loss
- Custom Metrics
- Custom Layers
- Custom Models

Object-Oriented Style

Custom Loss

```
def custom mean squared error(y true, y pred):
    return tf.math.reduce mean(tf.square(y true - y pred))
model.compile(optimizer=keras.optimizers.Adam(), loss=custom mean squared error)
class CustomMSE(keras.losses.Loss):
    def init (self, regularization factor=0.1, name="custom mse"):
        super(). init (name=name)
        self.regularization factor = regularization factor
    def call(self, y true, y pred):
        mse = tf.math.reduce mean(tf.square(y true - y pred))
        reg = tf.math.reduce mean(tf.square(0.5 - y pred))
        return mse + reg * self.regularization factor
model.compile(optimizer=keras.optimizers.Adam(), loss=CustomMSE())
```

Custom Metrics

```
class MyMetric(keras.metrics.Metric):
    def init (self, **kwargs):
        super(MyMetric, self). init (**kwargs)
        self.value = self.add weight(name="value", initializer="zeros")
        self.num = self.add weight(name="num", initializer="zeros")
    def update state(self, y true, y pred, sample weight=None):
        self.value.assign add(tf.constant(100.0))
        self.num.assign add(tf.constant(1.0))
    def result(self):
        return self.value
    def reset states(self):
        # The state of the metric will be reset at the start of each epoch.
        self.value.assign(0.0)
        self.num.assign(0.0)
model.compile(optimizer=keras.optimizers.Adam(), metric=MyMetric())
```

Custom Layers (1/2)

```
class ActivityRegularizationLayer(layers.Layer):
    def call(self, inputs):
        self.add loss(tf.reduce sum(inputs) * 0.1)
        return inputs # Pass-through layer.
inputs = keras.Input(shape=(784,), name="digits")
x = layers.Dense(64, activation="relu", name="dense 1")(inputs)
# Insert activity regularization as a layer
x = ActivityRegularizationLayer()(x)
x = layers.Dense(64, activation="relu", name="dense 2")(x)
outputs = layers. Dense (10, name = "predictions")(x)
model = keras.Model(inputs=inputs, outputs=outputs)
model.compile(
    optimizer=keras.optimizers.RMSprop(learning rate=1e-3),
    loss=keras.losses.SparseCategoricalCrossentropy(from logits=True),
```

Custom Layers (2/2)

```
class LogisticEndpoint(keras.layers.Layer):
   def init (self, name=None):
        super(LogisticEndpoint, self).__init__(name=name)
        self.loss fn = keras.losses.BinaryCrossentropy(from logits=True)
        self.accuracy fn = keras.metrics.BinaryAccuracy()
   def call(self, targets, logits, sample weights=None):
        # Compute the training-time loss value and add it
        # to the layer using `self.add loss()`.
        loss = self.loss fn(targets, logits, sample weights)
        self.add loss(loss)
        # Log accuracy as a metric and add it
        # to the layer using `self.add metric()`.
        acc = self.accuracy fn(targets, logits, sample weights)
        self.add metric(acc, name="accuracy")
        # Return the inference-time prediction tensor (for `.predict()`).
        return tf.nn.softmax(logits)
```